

corr_numericvars

March 1, 2025

0.0.1 This notebook performs and analyses correlation analysis. It first generates a correlation matrix and then visualizes the matrix. This is a precursor to Principle Component Analysis for consideration of dimensional reduction in preprocessing.

0.0.2 dependencies

```
[14]: import pandas as pd
import seaborn as sn
import matplotlib.pyplot as plt
import os
```

```
[48]: # makes relative paths work situationally ) I like access to the whole class at_
↳ one time)
os.chdir('C:/Users/entro/Desktop/Spring25_Semester/DataMining/SemesterProject/
↳ Bennett_CookePolitikos_CS368_StellarClassification/Scripts and notebooks')

print(os.getcwd())
```

C:\Users\entro\Desktop\Spring25_Semester\DataMining\SemesterProject\Bennett_CookePolitikos_CS368_StellarClassification\Scripts and notebooks

```
[76]: df = pd.read_csv('../Dataset/star_classification.csv')
df.head()
```

```
[76]:
```

	obj_ID	alpha	delta	u	g	r	\
0	1.237661e+18	135.689107	32.494632	23.87882	22.27530	20.39501	
1	1.237665e+18	144.826101	31.274185	24.77759	22.83188	22.58444	
2	1.237661e+18	142.188790	35.582444	25.26307	22.66389	20.60976	
3	1.237663e+18	338.741038	-0.402828	22.13682	23.77656	21.61162	
4	1.237680e+18	345.282593	21.183866	19.43718	17.58028	16.49747	

	i	z	run_ID	rerun_ID	cam_col	field_ID	spec_obj_ID	\
0	19.16573	18.79371	3606	301	2	79	6.543777e+18	
1	21.16812	21.61427	4518	301	5	119	1.176014e+19	
2	19.34857	18.94827	3606	301	2	120	5.152200e+18	
3	20.50454	19.25010	4192	301	3	214	1.030107e+19	

```
4  15.97711  15.54461    8102      301      3      137  6.891865e+18
```

```

      class  redshift  plate    MJD  fiber_ID
0  GALAXY  0.634794   5812  56354     171
1  GALAXY  0.779136  10445  58158     427
2  GALAXY  0.644195   4576  55592     299
3  GALAXY  0.932346   9149  58039     775
4  GALAXY  0.116123   6121  56187     842

```

```
[24]: df.dtypes
```

```

[24]: obj_ID      float64
      alpha      float64
      delta      float64
      u          float64
      g          float64
      r          float64
      i          float64
      z          float64
      run_ID      int64
      rerun_ID    int64
      cam_col     int64
      field_ID    int64
      spec_obj_ID float64
      class       object
      redshift    float64
      plate       int64
      MJD         int64
      fiber_ID    int64
      dtype: object

```

I would like to be sure that none of the IDs are nonarbitrary (such as representing a position) but for the time being this removes all IDs. It makes sense that some of these IDs are huge numbers, just a tendency of mine.

```

[ ]: # creates list of attributes with float64 data types/ ints are indexes
num_cols = df.columns[df.dtypes == "float64"]
print(f"Numeric variables: {num_cols} \n\n")

# creates pandas data frame that subsets imported df to include only floating_
↳ point values
numvar_df = df[num_cols]
print(numvar_df.head())

# removes any additional columns containing ID values
# "inplace" param indicates that it executes the change to the df that it is_
↳ called on
# as opposed to returning a df that can be saved to a new variable

```

```
numvar_df.drop(list(numvar_df.filter(regex = 'ID')), axis = 1, inplace = True)

print("\n\n")
numvar_df.head()
```

```
Numeric variables: Index(['obj_ID', 'alpha', 'delta', 'u', 'g', 'r', 'i', 'z',
                          'spec_obj_ID',
                          'redshift'],
                          dtype='object')
```

	obj_ID	alpha	delta	u	g	r	\
0	1.237661e+18	135.689107	32.494632	23.87882	22.27530	20.39501	
1	1.237665e+18	144.826101	31.274185	24.77759	22.83188	22.58444	
2	1.237661e+18	142.188790	35.582444	25.26307	22.66389	20.60976	
3	1.237663e+18	338.741038	-0.402828	22.13682	23.77656	21.61162	
4	1.237680e+18	345.282593	21.183866	19.43718	17.58028	16.49747	

	i	z	spec_obj_ID	redshift
0	19.16573	18.79371	6.543777e+18	0.634794
1	21.16812	21.61427	1.176014e+19	0.779136
2	19.34857	18.94827	5.152200e+18	0.644195
3	20.50454	19.25010	1.030107e+19	0.932346
4	15.97711	15.54461	6.891865e+18	0.116123

C:\Users\entro\AppData\Local\Temp\ipykernel_16808\1746605293.py:12:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
numvar_df.drop(list(numvar_df.filter(regex = 'ID')), axis = 1, inplace = True)
```

```
[ ]:      alpha      delta      u      g      r      i      z  \
0  135.689107  32.494632  23.87882  22.27530  20.39501  19.16573  18.79371
1  144.826101  31.274185  24.77759  22.83188  22.58444  21.16812  21.61427
2  142.188790  35.582444  25.26307  22.66389  20.60976  19.34857  18.94827
3  338.741038 -0.402828  22.13682  23.77656  21.61162  20.50454  19.25010
4  345.282593  21.183866  19.43718  17.58028  16.49747  15.97711  15.54461

      redshift
0  0.634794
1  0.779136
2  0.644195
3  0.932346
```

4 0.116123

Creates a correlation matrix on the numeric data

```
[ ]: corr_matrix = numvar_df.corr()

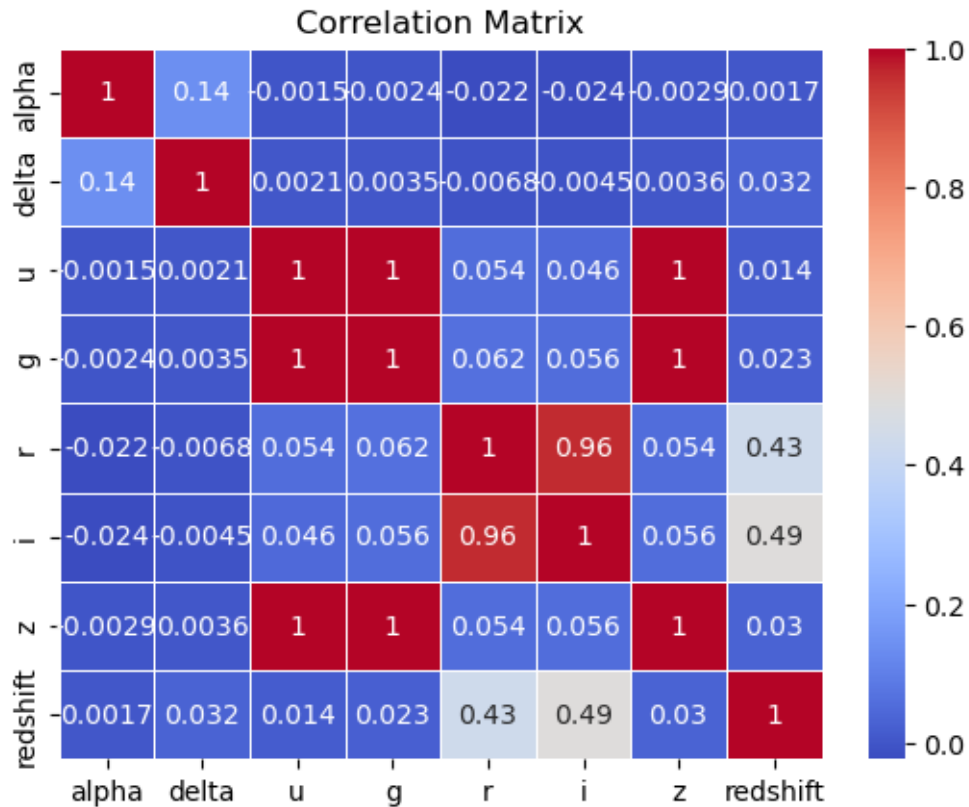
corr_matrix.style\
    .format(precision = 3)\
    .format_index(str.upper, axis = 0)\
    .format_index(str.upper, axis = 1)\
    .background_gradient(cmap='coolwarm')
```

```
[ ]: <pandas.io.formats.style.Styler at 0x2371f32aae0>
```

other than a screen shot I haven't figured out how to export the above heatmap

```
[ ]: sn.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()

# saves visual to IMGs folder and returns to current directory
current = os.getcwd()
os.chdir('C:/Users/entro/Desktop/Spring25_Semester/DataMining/SemesterProject/
↳Bennett_CookePolitikos_CS368_StellarClassification/IMGs/Tables')
plt.savefig("correlation.svg", format="svg")
os.chdir(current)
```



<Figure size 640x480 with 0 Axes>

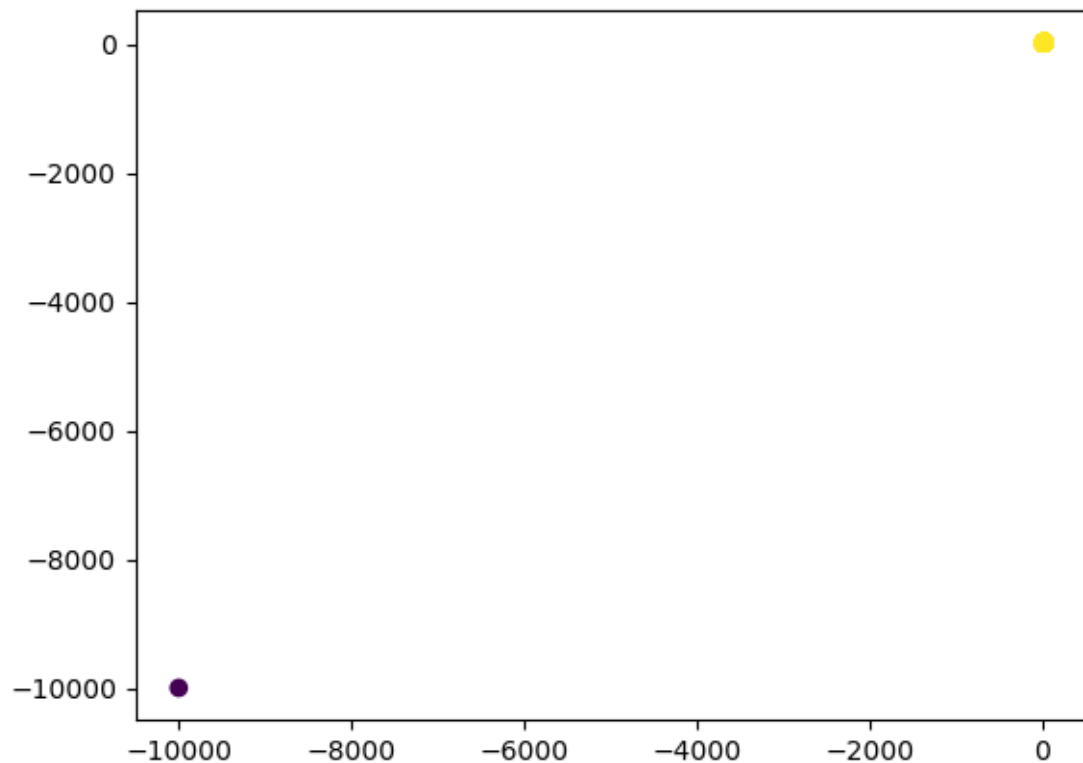
The above looks like ass.

Needs: * consistent sigfigs * better label orientation (rotate) * better labels for single letters (for communication purposes)

Analysis:

- low correlation allows for much of the data to effectively inform on the model which is great
- high correlation:
 - u and g
 - i and r
 - z and both u & g (because they are highly correlated with each other, makes sense that if z is strongly correlated with one it must be with the other)

```
[56]: x = numvar_df['z']
      y = numvar_df['u']
      z = numvar_df['g']
      plt.scatter(x = x, y = y, c = z)
      plt.show()
```



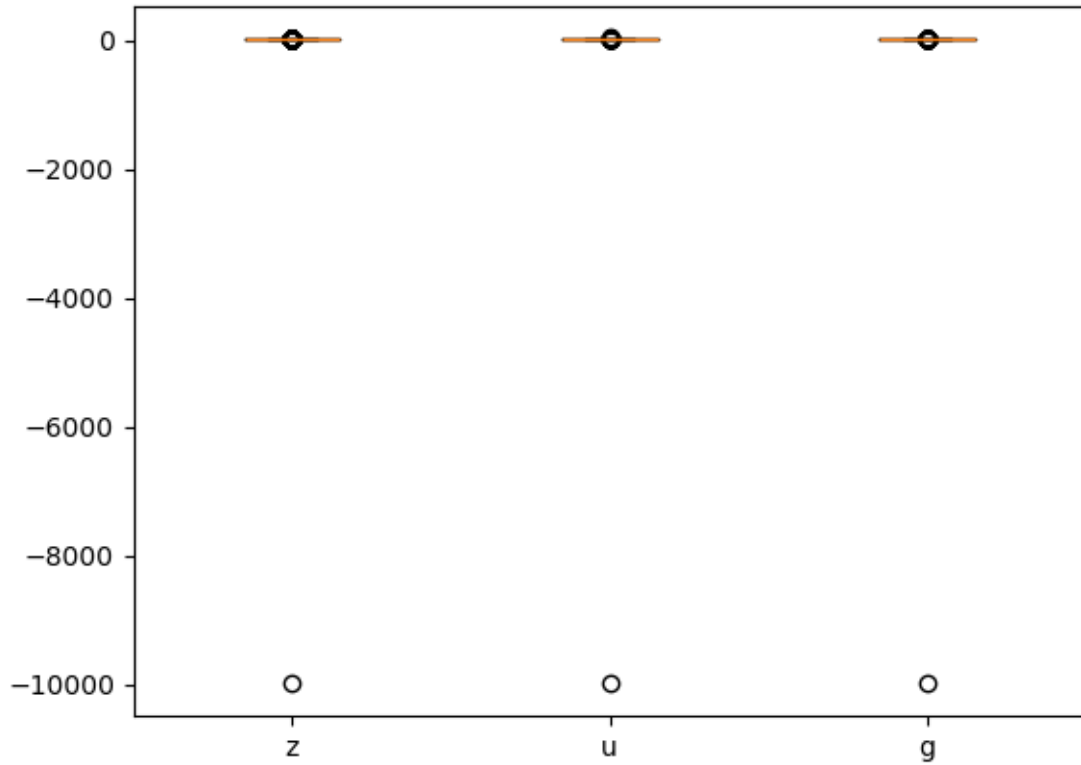
Above scatter plot is meant to show how the distributions of g and u vary respective to z but it looks like there is some outliers that are extending the range of the data which is confirmed in the boxplot below.

```
[ ]: plt.boxplot([numvar_df['z'], numvar_df['u'], numvar_df['g']], labels=
    ↪ ['z', 'u', 'g'])
plt.show()
```

C:\Users\entro\AppData\Local\Temp\ipykernel_16808\2871417091.py:1:

MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.

```
plt.boxplot([numvar_df['z'], numvar_df['u'], numvar_df['g']], labels=
['z', 'u', 'g'])
```



The box plot below removes all values below zero (the outlier seems to be a single object with a nonsensical value for z,u,and g though they could be different objects in theory)

```
[68]: numvar_df[numvar_df['z']>= 0]
```

```
[68]:
```

	alpha	delta	u	g	r	i \
0	135.689107	32.494632	23.87882	22.27530	20.39501	19.16573
1	144.826101	31.274185	24.77759	22.83188	22.58444	21.16812
2	142.188790	35.582444	25.26307	22.66389	20.60976	19.34857
3	338.741038	-0.402828	22.13682	23.77656	21.61162	20.50454
4	345.282593	21.183866	19.43718	17.58028	16.49747	15.97711
...
99995	39.620709	-2.594074	22.16759	22.97586	21.90404	21.30548
99996	29.493819	19.798874	22.69118	22.38628	20.45003	19.75759
99997	224.587407	15.700707	21.16916	19.26997	18.20428	17.69034
99998	212.268621	46.660365	25.35039	21.63757	19.91386	19.07254
99999	196.896053	49.464643	22.62171	21.79745	20.60115	20.00959

	z	redshift
0	18.79371	0.634794
1	21.61427	0.779136
2	18.94827	0.644195

```

3      19.25010  0.932346
4      15.54461  0.116123
...
99995  20.73569  0.000000
99996  19.41526  0.404895
99997  17.35221  0.143366
99998  18.62482  0.455040
99999  19.28075  0.542944

```

```
[99999 rows] x 8 columns]
```

```
[70]: plt.boxplot([numvar_df['z'][numvar_df['z']>= 0],
    ↪numvar_df['u'][numvar_df['u']>= 0], numvar_df['g'][numvar_df['g']>= 0]],
    ↪labels= ['z','u','g'])
plt.show()
```

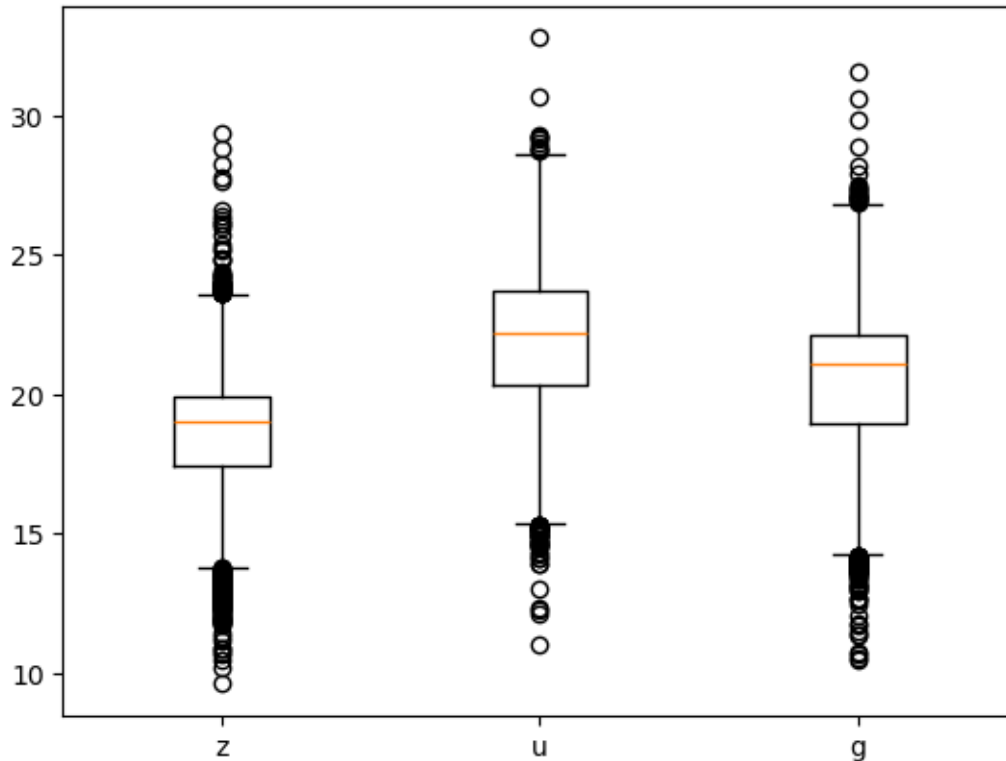
C:\Users\entro\AppData\Local\Temp\ipykernel_16808\3115795730.py:1:

MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.

```

plt.boxplot([numvar_df['z'][numvar_df['z']>= 0],
numvar_df['u'][numvar_df['u']>= 0], numvar_df['g'][numvar_df['g']>= 0]], labels=
['z','u','g'])

```



The above correction shows distributions within the same range of values. This means that with the nonsensical value removed, the distributions should not need to be normalized to visualize effectively (though normalization might be valuable in a larger context). The plot below is the same scatter plot with the nonsense value removed.

```
[74]: z = numvar_df['z'][numvar_df['z']>= 0]
      u = numvar_df['u'][numvar_df['u']>= 0]
      g = numvar_df['g'][numvar_df['g']>= 0]

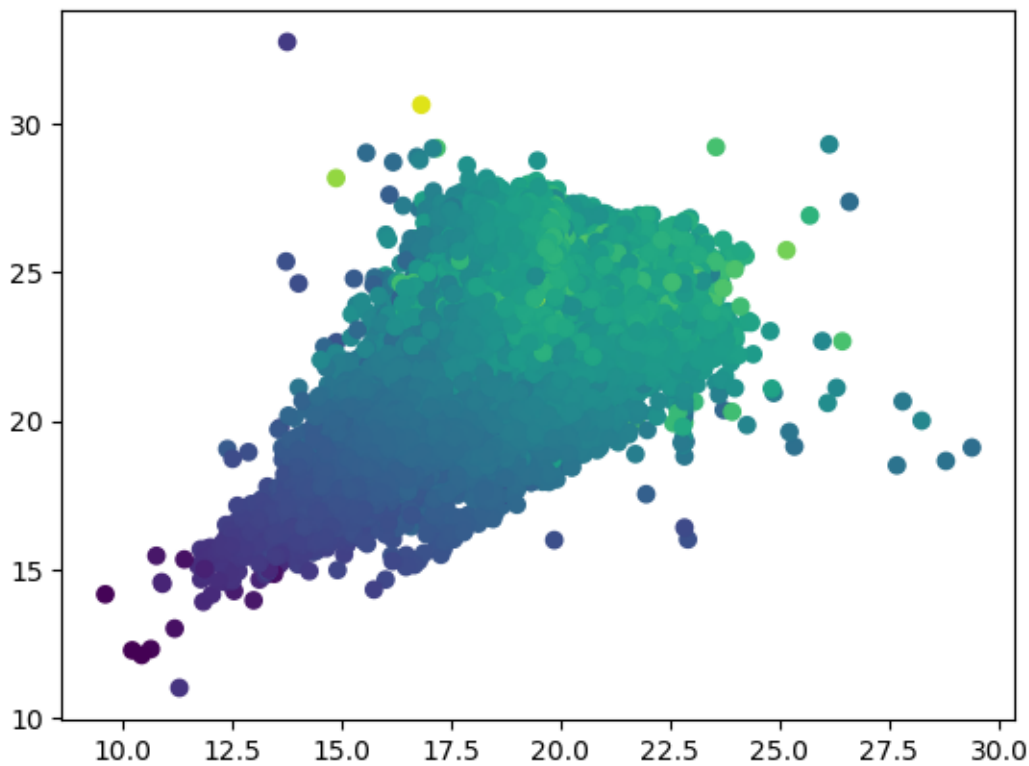
      lengths = f"z length: {len(z)}, u length: {len(u)}, g length {len(g)}"

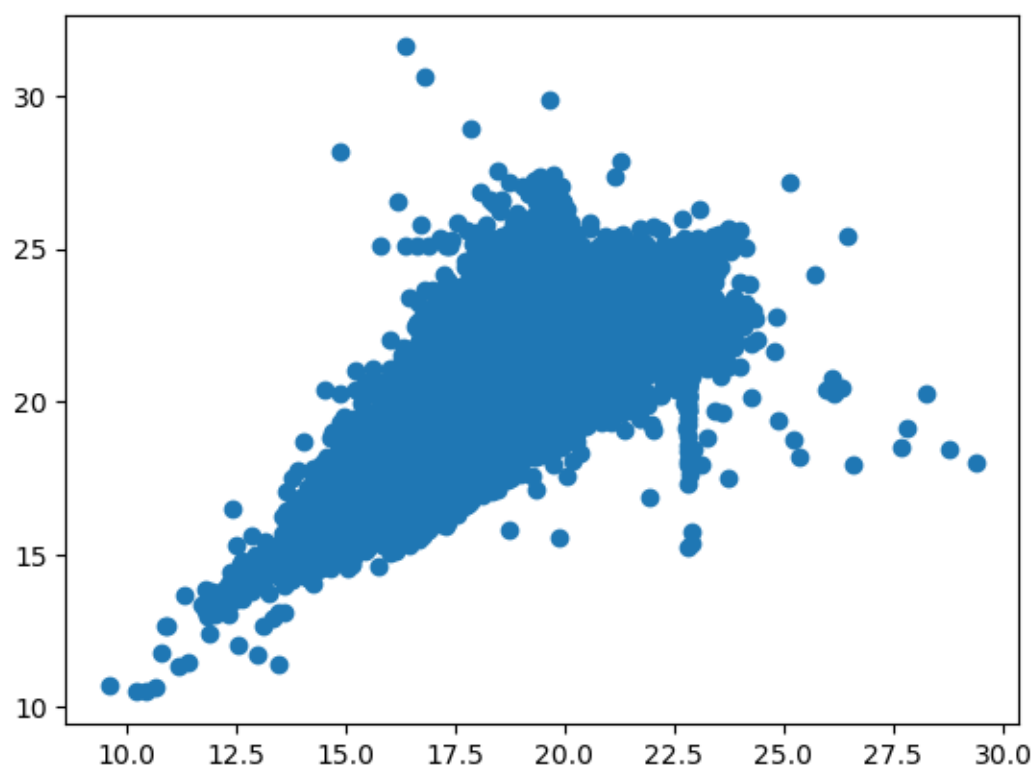
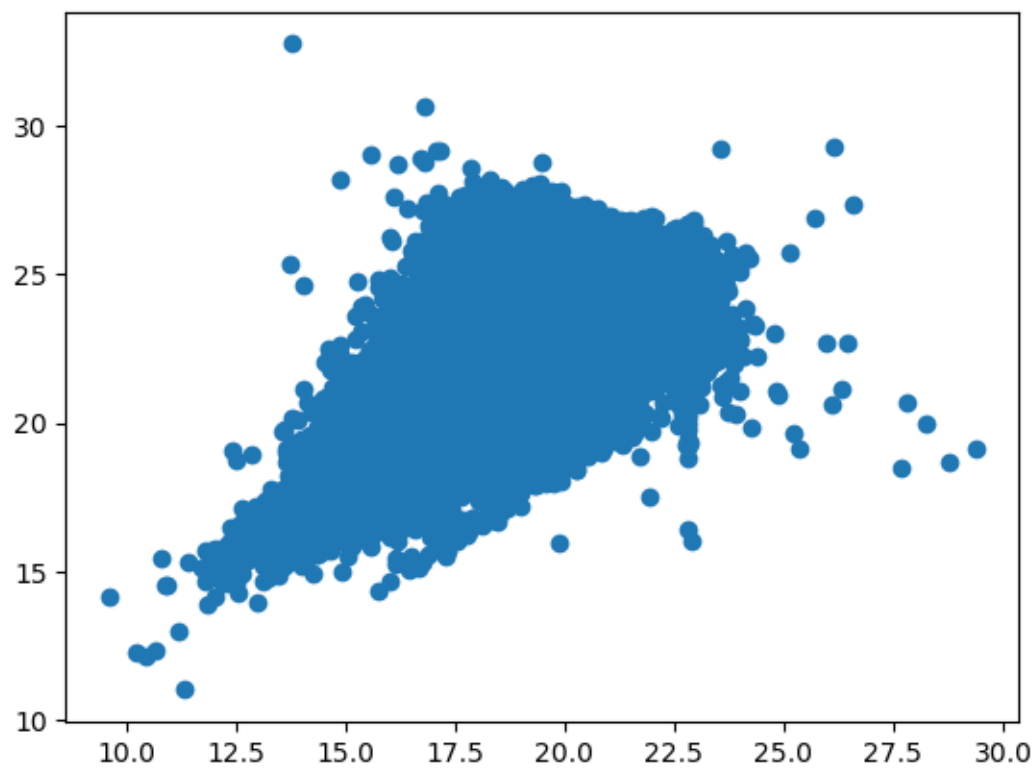
      plt.scatter(x = z, y = u, c = g)
      plt.show()

      plt.scatter(x = z, y = u)
      plt.show()

      plt.scatter(x = z, y = g)
      plt.show()

      print(lengths)
```





z length: 99999, u length: 99999, g length 99999

determines the location of the nonsensical values

```
[77]: u_loc = df['u'].idxmin()
      g_loc = df['g'].idxmin()
      z_loc = df['z'].idxmin()

      if z_loc == u_loc & z_loc == g_loc:
          print(f'the location of the object is at row {u_loc}')
```

the location of the object is at row 79543

Review: This notebook provided information about what objects correlate with each other. This is valuable in determining what attribute will make redundant contributions to our eventual models. It also performed some initial outlier analysis but did not create a threshold for making the determination of what values will be considered outliers and only identified one outlier. Relative to the second boxplot, there may be other outlier candidates. But it might be good to consider this across all of the attributes of the data objects.